

## Lower Bounds

- So far we have concentrated on the question: *Given a problem  $\Pi$  can we construct an algorithm that solves  $\Pi$  in  $O(f(n))$  time?* In other words, is there an algorithm  $A$  that satisfies:

$$T_A(n) = \max_{|X|=n} T_A(X) \in O(f(n))$$

- *Goal:* find  $f(n)$  that grows as *slowly* as possible
- Today we concentrate on proving statements of the form: *any algorithm that solves  $X$  must take  $\Omega(g(n))$  time*

$$T_{\Pi}(n) = \min_{A \text{ solves } \Pi} T_A(n) = \min_{A \text{ solves } \Pi} \max_{|X|=n} T_A(X) \in \Omega(g(n))$$

- The complexity  $T_{\Pi}(n)$  of a problem  $\Pi$  is the complexity of the *best algorithm* that solves  $\Pi$
- *New Goal:* find  $g(n)$  that grows as *fast* as possible

## Lower Bounds...

- The ratio of the slowest growing upper bound to the fastest growing lower bound,  $f(n)/g(n)$ , is a kind of “gap” for  $\Pi$

Example: If  $\Pi$  has a lower bound of  $\Omega(n \log n)$  and the best known algorithm solves  $\Pi$  in  $O(n \log^2 n)$  time then there is a gap of  $\log n$  for  $\Pi$

- When  $A$  satisfies  $f(n) \in \Theta(g(n))$ ,  $A$  is *optimal*
- Improving lower bounds is considerably harder than improving upper bounds, because a lower bound applies to *all* algorithms that solve  $\Pi$
- *Problem:* what do we mean by ‘*all algorithms*’?

## Models of Computation

- We first specify the *model of computation*, i.e., the kinds of algorithms allowed and the cost of the model operations
- Lower bounds apply to a specific model of computation
- Model today: *decision tree*, a  $k$ -ary tree such that:
  - Each internal node is labeled with a query about the input
  - Edges out of a node correspond to answer to the query
  - Each leaf is labeled with a possible output
  - A specific computation is a path from root to a leaf, where the answers tell us what to compute next
  - Cost is the number of queries asked
- Each query has  $\leq k$  branches,  $N$  possible total outputs  $\Rightarrow \geq \lceil \log_k N \rceil$  queries are needed in the worst case

3

## Lower Bounds with Decision Trees

- Lower bounds to many problems can be obtained using decision trees
  - Searching a sorted list
  - Searching an unsorted list
  - Finding the  $i$ -th smallest element of a list
  - Finding the mode of a list
  - Sorting a list
  - Merging two sorted lists
  - Find the elements common to two lists
  - Determine if all elements of a list are distinct
  - Determine if two lists have the same elements

4

## Comparison Trees

- A *comparison tree* is a type of decision tree where each query involves the comparison of two values
  - Each internal node  $v$  is labeled with a comparison  $x: y$  for some input keys  $x$  and  $y$
  - Each internal node has 2 or 3 outcomes
  - Each leaf is labeled with an output of  $\Pi$  on some input of size  $n$
  - For each input  $x$ , there is a  $\text{path}(x)$  from root to a leaf such that every edge  $(u, v)$  in  $\text{path}(x)$  is labeled with the comparison performed at  $u$
  - Tree is correct if for every  $x$ , leaf in  $\text{path}(x)$  is a valid output

5

## Comparison Trees...


- You have one tree for each combination of input size and algorithm for problem  $\Pi$
- The number of leaves in a decision tree of order  $n$  for  $\Pi$  is greater than or equal to the number of distinct outputs of  $\Pi$  on inputs of size  $n$
- A  $k$ -ary tree with  $L$  leaves has height at least  $\lceil \log_k L \rceil$ , a lower bound for problem  $\Pi$

6

## Example: Searching a Sorted List

- Each searching strategy and input size has a tree
- Below is the tree for standard *binary search*

2	3	5	7	11	13	17	19
---	---	---	---	----	----	----	----



- Depth is in  $\Omega(\log n)$  Independent of tree organization

7

## Exercise

1. Describe the decision tree for the following search strategies:
  - Linear search
  - Jump search
  - Exponential search
2. Use comparison trees to derive a lower bound on the complexity of merging two sorted lists of size  $n$  each

8

## Lower Bounds for Sorting

- How fast can we sort?
  - Answer depends on computational model
- So far:
  - Insertion sort takes  $\Theta(n^2)$  (worst case)
  - Quicksort takes  $\Theta(n \log n)$  (expected)
  - Merge sort takes  $\Theta(n \log n)$  (worst case)
  - Heapsort takes  $\Theta(n \log n)$  (worst case)
- Can we do better than  $\Theta(n \log n)$ ?
- These algorithms share the same model
  - Will provide a lower bound for this model and then beat it (for *restricted* inputs) by changing the model

## Comparison Sort Model

- Uses the comparison tree model, i.e., the basic operation is the *comparison of two elements*
  - Only use comparisons to determine relative order (resulting algorithm is called a *comparison sort*)
  - Only count comparisons to determine complexity
- Lower bounds
  - $\Omega(n)$  to examine all the input
  - All sorts seen so far are comparison sorts and take  $\Omega(n \log n)$  in the worst case
  - Will show  $\Omega(n \log n)$  lower bound for this model

2

# Comparison Trees for Sorting

- A *comparison tree* is used as an abstraction of a comparison sort
- The tree represents the set of *all* possible comparisons made by a fixed algorithm on inputs of a fixed size
- Abstracts away everything else, such as control and data movement
- Only comparisons are counted


11

## Example

- Insertion sort,  $n = 3$ , e.g., sort  $\langle 3, 5, 2 \rangle$

```

Sort( $A, n$ )
1. for  $j \leftarrow 2$  to  $n$  do
2.    $k \leftarrow A[j]$ 
3.    $i \leftarrow j - 1$ 
4.   while  $i > 0$  and  $A[i] > k$  do
5.      $A[i+1] \leftarrow A[i]$ 
6.      $i \leftarrow i - 1$ 
7.    $A[i+1] \leftarrow k$ 
    
```



- Each node labeled with *original* element indices
- Each leaf labeled by permutation found by algorithm
- Path in bold corresponds to  $a_3 \leq a_1 \leq a_2$

12

## More generally...

- Want to sort  $\langle a_1, \dots, a_n \rangle$
- Each internal node has label  $a_i : a_j$ , where  $i, j \in \{1, \dots, n\}$
- Left subtree contains comparisons performed *after* determining that  $a_i \leq a_j$
- Right subtree contains subsequent comparisons for the case  $a_i > a_j$
- Each leaf node has a permutation of  $\langle 1, \dots, n \rangle$  that corresponds to the correct sorted order of the input

13

## Properties

For a particular (deterministic) algorithm

- One tree for each  $n$
- All possible execution traces are represented
- A specific run  $\Rightarrow$  a path from root to leaf
- How many leaves does a decision tree have?
- What is the length of longest path from root to leaf?  
Depends on the algorithm!
  - insertion sort?
  - heapsort?
  - merge sort?
  - quicksort?

14

## Lower Bound for Sorting

**Theorem.** Any decision tree for sorting  $n$  elements has height  $\Omega(n \log n)$

**Proof.**

A binary tree of height  $h$  has  $\leq 2^h$  leaves

Decision tree for correct algorithm has  $\geq n!$  leaves

$\Rightarrow n! \leq \# \text{leaves} \leq 2^h \Rightarrow h \geq \log n!$


**Corollary.** Heapsort and Merge sort are asymptotically optimal (under the comparison model of sorting)

15

## Other Lower Bounds

- Given a problem  $X$ , one can often use a **reduction** from *Sorting* to  $X$  to show that  $X$  has a  $\Omega(n \log n)$  lower bound as well
- The reduction must require  $\omega(n \log n)$  time

*Example:*  $\text{Sorting} \leq_{O(n)} \text{Convex-Hull}$



*Example:*  $\text{Convex Hull} \leq_{O(n)} \text{Triangulation}$

16



## Digit-Based Sorting

- So far, when sorting, we have viewed the input keys as abstract objects that can only be examined via comparisons
- Now, we view each input key as a sequence of “digits”
- Digits can be individually manipulated
- New point of view leads to:
  - Fast sorting algorithm (radix sort)
  - Online data structure (tries)

17

## Sorting in Linear Time

- Non-comparison sorts
- Need additional assumptions about items to be sorted

### Example: Counting Sort

*Assumption*: input integers in  $\{0, 1, 2, \dots, k\}$

– Input:  $A[1..n]$     Output:  $B[1..n]$     Auxiliary:  $C[0..k]$

– Idea: for each  $i$  in  $1..n$  compute rank of  $A[i]$

–  $C[i] = \text{rank of } A[i] = \# \text{ elements from } A \text{ that are } \leq i$

18

```

COUNTING-SORT( $A, B, n, k$ )
  for  $i \leftarrow 0$  to  $k$ 
    do  $C[i] \leftarrow 0$ 
  for  $j \leftarrow 1$  to  $n$ 
    do  $C[A[j]] \leftarrow C[A[j]] + 1$ 
  for  $i \leftarrow 1$  to  $k$ 
    do  $C[i] \leftarrow C[i] + C[i - 1]$ 
  for  $j \leftarrow n$  downto 1
    do  $B[C[A[j]]] \leftarrow A[j]$ 
        $C[A[j]] \leftarrow C[A[j]] - 1$ 


```

**Analysis.**

- Running time:  $\Theta(n + k)$  which is  $\Theta(n)$  if  $k = O(n)$
- How big a  $k$  is practical?  
32-bit numbers? 16-bit? 8-bit?

19

## Example




20

## What if $k$ is big?

- Assume each input number  $a$  has  $d$  “digits”

$$a = a_d a_{d-1} \dots a_2 a_1$$

- IBM’s card sorting machine



- Two strategies based on sorting one column at a time
  - most significant digit first
  - least significant digit first

21


## Radix Sort

Input value  $x$  consists of  $d$  digits:  $x = x_d x_{d-1} \dots x_2 x_1$

**RADIX-SORT**( $A, d$ )

**for**  $i \leftarrow 1$  **to**  $d$

**do** use a stable sort to sort array  $A$  on digit  $i$



22

## Analysis

- Correctness
  - Prove by induction on  $d$  (number of passes)
    - Assume input already sorted on digits  $1, \dots, i-1$
    - Argue that sort on digit  $i$ , leaves digits  $1, \dots, i$  sorted
- Time complexity:
  - $O(d(n+k))$  when coupled with counting sort
  - depends on
    - stable sort used
    - $d$ , which depends on  $n$  and number base used

23

## How do you break keys into digits?

- Each key has  $b$  bits
- Each digit is  $r$  bits  $\Rightarrow d = \lceil b/r \rceil$
- With counting sort,  $k = 2^r - 1$ 
  - Example:  $b=32, r=8 \Rightarrow d=32/8=4, k=255$
- Time:  $\Theta(b/r (n + 2^r))$
- How do you choose  $r$ ?
  - $r \approx \log n \Rightarrow$  time is  $\Theta(bn/\log n)$
  - $r < \log n \Rightarrow b/r > b/\log n$  but  $(n+2^r) = \Omega(n)$
  - $r > \log n \Rightarrow (n+2^r)$  gets big quickly

24

## Merge Sort or Radix Sort?

- Sort 1 million 32-bit integers
- Merge sort performs 20 “passes”
- Since  $\lceil \log 10^6 \rceil = 20$ , radix sort performs  $\lceil 32/20 \rceil = 2$  calls to counting sort
- Each call to counting sort requires 4 passes

25

## Bucket Sort

- Input: array  $A[1..n]$  of numbers in  $[0,1)$   
Uses auxiliary array  $B[0..n-1]$  of linked lists
- Algorithm
  - Divide  $[0,1)$  into  $n$  equal size buckets
  - Place each input value into corresponding bucket
  - Sort the buckets independently
  - Concatenate
- Works well if input is uniformly distributed

26

BUCKET-SORT( $A, n$ )

**for**  $i \leftarrow 1$  **to**  $n$


**do** insert  $A[i]$  into list  $B[\lfloor n \cdot A[i] \rfloor]$

**for**  $i \leftarrow 0$  **to**  $n - 1$

**do** sort list  $B[i]$  with insertion sort

concatenate lists  $B[0], B[1], \dots, B[n - 1]$  together in order

**return** the concatenated lists



27

## Correctness

Consider to arbitrary keys:  $A[i]$  and  $A[j]$

$$A[i] < A[j] \Rightarrow \lfloor n \cdot A[i] \rfloor \leq \lfloor n \cdot A[j] \rfloor$$

$\Rightarrow A[i]$  is placed on same bucket as  $A[j]$

**or** in bucket with lower index

If same bucket, then insertion sort fixes the order

If different bucket, concatenation fixes the order

28

## Probabilistic Analysis

- Assume input generated by random process
- Let  $n_i$  denote size of  $B[i]$  (a random variable)

$$T(n) = \Theta(n) + \sum_{i=0}^{n-1} O(n_i^2)$$

$$E[T(n)] = \Theta(n) + \sum_{i=0}^{n-1} O(E[n_i^2])$$

**Claim.**  $E[(n_i)^2] = 2 - 1/n$

29

## Proof

$$X_{ij} = I\{A[j] \text{ falls in bucket } i\}$$

$$n_i = \sum_{j=1}^n X_{ij}$$

$$E[n_i^2] = E\left[\left(\sum_{j=1}^n X_{ij}\right)^2\right]$$

$$= E\left[\sum_{j=1}^n X_{ij}^2 + 2\sum_{j=1}^{n-1} \sum_{k=j+1}^n X_{ij} X_{ik}\right]$$

$$= \sum_{j=1}^n E[X_{ij}^2] + 2\sum_{j=1}^{n-1} \sum_{k=j+1}^n E[X_{ij} X_{ik}]$$

30

### Key observations

- $X_{ij}^2$  is a decision variable, same as  $X_{ij}$
- $X_{ij}$  and  $X_{ik}$  are independent random variables

$$\begin{aligned} & \sum_{j=1}^n E[X_{ij}^2] + 2 \sum_{j=1}^{n-1} \sum_{k=j+1}^n E[X_{ij} X_{ik}] \\ &= \sum_{j=1}^n E[X_{ij}^2] + 2 \sum_{j=1}^{n-1} \sum_{k=j+1}^n E[X_{ij}] E[X_{ik}] \\ &= \sum_{j=1}^n \frac{1}{n} + 2 \sum_{j=1}^{n-1} \sum_{k=j+1}^n \frac{1}{n} \cdot \frac{1}{n} = 1 + 2 \binom{n}{2} \frac{1}{n^2} \\ &= 1 + 2 \frac{n(n-1)}{2} \cdot \frac{1}{n^2} = 1 + \frac{n-1}{n} = 2 - \frac{1}{n} \quad \blacksquare \end{aligned}$$

31