


Search Trees

- Support dynamic set operations:
 - Insert - Min - Predecessor
 - Delete - Max - Successor
 - Search
- Can simultaneously be used as a *dictionary* and as a *priority queue*
- Operations run in $O(h)$ time, h = height of tree
- Many variants: binary search trees, red-black trees, AVL trees, 2-3-4 trees, B-trees, randomized search trees, splay trees, etc.

1


Binary Search Trees

- A type of search tree
- *Linked* binary tree structure
- Each node p contains:
 - *key*: unique value from totally ordered domain
 - *left*: points to left child
 - *right*: points to right child
 - *p*: points to parent
- $root[T]$ points to root node (also denoted by $T.root$)
- Binary search tree property:
 - if q is a descendant of $p.left \Rightarrow q.key < p.key$
 - if q is a descendant of $p.right \Rightarrow q.key > p.key$




2

Examples, $n = 6$



- How many BST's are there for each n ?



3

Traversals

- Systematic way to visit all nodes
- Three types: *in-order*, *pre-order*, *post-order*
- Each takes $\Theta(n)$ time


INORDER-TREE-WALK(x)

if $x \neq \text{NIL}$

 then INORDER-TREE-WALK($\text{left}[x]$)

 print $\text{key}[x]$

 INORDER-TREE-WALK($\text{right}[x]$)



4

Search

- Initial call $\text{Tree-Search}(\text{root}[T], k)$

```


TREE-SEARCH( $x, k$ )
  if  $x = \text{NIL}$  or  $k = \text{key}[x]$ 
    then return  $x$ 
  if  $k < \text{key}[x]$ 
    then return TREE-SEARCH( $\text{left}[x], k$ )
  else return TREE-SEARCH( $\text{right}[x], k$ )

```

- Time? $O(h)$

5

Minimum and Maximum



- Can be found using tree structure alone
- Follow rightmost path to max element
- Follow leftmost path to min element

6

Minimum and Maximum

<pre> TREE-MINIMUM(<i>x</i>) while <i>left</i>[<i>x</i>] ≠ NIL do <i>x</i> ← <i>left</i>[<i>x</i>] return <i>x</i> </pre>	<pre> TREE-MAXIMUM(<i>x</i>) while <i>right</i>[<i>x</i>] ≠ NIL do <i>x</i> ← <i>right</i>[<i>x</i>] return <i>x</i> </pre>
---	---

- Can also write recursively.
- Time? $O(h)$

7

Successor

- Successor of x is node y containing the smallest key $> x.key$
- Can be found using tree structure (no key comparisons necessary!)
- Two cases depending on whether $x.right = \text{NIL}$

```

TREE-SUCCESSOR(x)
  if right[x] ≠ NIL
    then return TREE-MINIMUM(right[x])
  y ← p[x]
  while y ≠ NIL and x = right[y]
    do x ← y
    y ← p[y]
  return y

```



8

Insertion

- z initialized with NIL left and right pointers

TREE-INSERT(T, z)

$y \leftarrow \text{NIL}$

$x \leftarrow \text{root}[T]$

while $x \neq \text{NIL}$

do $y \leftarrow x$

if $\text{key}[z] < \text{key}[x]$

then $x \leftarrow \text{left}[x]$

else $x \leftarrow \text{right}[x]$

$p[z] \leftarrow y$

if $y = \text{NIL}$


then $\text{root}[T] \leftarrow z$

▷ Tree T was empty

else if $\text{key}[z] < \text{key}[y]$

then $\text{left}[y] \leftarrow z$

else $\text{right}[y] \leftarrow z$



9

Tree-Delete(T, z)

TREE-DELETE is broken into three cases.

Case 1: z has no children.


- Delete z by making the parent of z point to NIL, instead of to z .

Case 2: z has one child.

- Delete z by making the parent of z point to z 's child, instead of to z .


Case 3: z has two children.

- z 's successor y has either no children or one child. (y is the minimum node—with no left child—in z 's right subtree.)
- Delete y from the tree (via Case 1 or 2).
- Replace z 's key and satellite data with y 's.



10

Examples



11

Analysis

- All operations take $O(h)$ time and h is $\Theta(\log n)$ in the best case and $\Theta(n)$ in the worst case
- Key to efficiency is to restructure the tree when h gets too big
- Does randomness help?


12

Randomly Built BST

- Given a set of n distinct keys enter them in random order into empty BST
- Each of $n!$ permutations equally likely for insertion order
- Different from assuming that every BST on n keys is equally likely.
- Can show:
 - average node depth is $O(\log n)$
 - expected height is $O(\log n)$

13

Example: $n = 500$



14

Using a BST to Sort

BST-Sort(A)

$T = \emptyset$


for $i = 1$ **to** n **do**

 Tree-Insert($T, A[i]$)

 In-order(T)

end

- Example: sort $\langle 4, 1, 8, 6, 3, 7, 5 \rangle$
- Time?



15

Analysis

- The in-order traversal takes $O(n)$ time.
- How long do n tree inserts take?

$$P_n = \sum_{v \in T} \text{depth}(v)$$

- Worst case
- Best case
- Average case
- Running time is always $\Omega(n \log n)$ and $O(n^2)$
 $O(\log n)$ height $\Rightarrow O(n \log n)$ running time

16

Average Case

Let T be a BST of n nodes

Internal path length $P_n := \sum_{x \in T} \text{depth}(x)$

$$P_1 = 0 \text{ and } P_n = P_{i-1} + P_{n-i} + n - 1$$


Define $P_0 := 0$ and let $\bar{P}_n := E(P_n)$

Then, if root has rank $R_n = i$

$$P_n = P_{i-1} + P_{n-i} + n - 1, \text{ and}$$

$$\bar{P}_n = \frac{1}{n} \sum_{i=1}^n (\bar{P}_{i-1} + \bar{P}_{n-i} + n - 1)$$

$$n\bar{P}_n = 2 \sum_{i=0}^{n-1} \bar{P}_i + n(n-1)$$



17

Average Case...

$$n\bar{P}_n = 2 \sum_{i=0}^{n-1} \bar{P}_i + n(n-1)$$

$$(n-1)\bar{P}_{n-1} = 2 \sum_{i=0}^{n-2} \bar{P}_i + (n-1)(n-2)$$

$$n\bar{P}_n - (n-1)\bar{P}_{n-1} = 2\bar{P}_{n-1} + 2(n-1)$$

$$n\bar{P}_n = (n+1)\bar{P}_{n-1} + 2(n-1)$$

$$n\bar{P}_n \leq (n+1)\bar{P}_{n-1} + 2n$$

$$\frac{\bar{P}_n}{n+1} \leq \frac{\bar{P}_{n-1}}{n} + \frac{2}{n+1}$$

18

Average Case...

$$\begin{aligned}
 \frac{\bar{P}_n}{n+1} &\leq \frac{\bar{P}_{n-1}}{n} + \frac{2}{n+1} \\
 &\leq \frac{\bar{P}_{n-2}}{n-1} + \frac{2}{n} + \frac{2}{n+1} \\
 &\leq \frac{\bar{P}_{n-3}}{n-2} + \frac{2}{n-1} + \frac{2}{n} + \frac{2}{n+1} \\
 &\vdots \\
 &\leq \frac{\bar{P}_1}{2} + \frac{2}{3} + \frac{2}{4} + \cdots + \frac{2}{n} + \frac{2}{n+1} < 2H_{n+1}
 \end{aligned}$$

$$\Rightarrow \bar{P}_n \leq 2(n+1)H_{n+1} = O(n \log n)$$

19

Analysis

- How long do n tree inserts take?

$$P_n = \sum_{v \in T} \text{depth}(v)$$

- Worst case is $\Theta(n^2)$.
- Best case is $\Theta(n \log n)$
- Average case (when $n!$ permutations equally likely) is $\Theta(n \log n)$
- What algorithm does this remind you of?

20

Relation to Quicksort

- BST sort and quicksort make the same comparisons, although in different order
Example: $\langle 4, 1, 8, 6, 3, 7, 5 \rangle \Leftrightarrow \langle 5, 7, 3, 6, 8, 1, 4 \rangle$
- Worst case is $\Theta(n^2)$
- Best case is $\Theta(n \log n)$
- Average case comes from Quicksort analysis
 \Rightarrow Use Randomized BST Sort
 1. Randomly permute A
 2. BST-Sort(A)

21

Expected Tree Height

Let X_n = height of T (a random variable)

$Y = 2^{X_n}$, exponential height of T (also a r.v.)

Suppose the root has rank i , $1 \leq i \leq n$. Then,

$$X_n = 1 + \max\{X_{i-1}, X_{n-i}\}$$

$$Y_n = 2^{X_n} = 2 \cdot \max\{Y_{i-1}, Y_{n-i}\} \leq 2 \cdot (Y_{i-1} + Y_{n-i})$$

$$E[Y_n] \leq \frac{1}{n} \sum_{i=1}^n 2 \cdot E(Y_{i-1} + Y_{n-i})$$


$$= \frac{4}{n} \sum_{i=1}^n E[Y_{i-1}] = \frac{4}{n} \sum_{k=0}^{n-1} E[Y_k]$$

22

Claim. $E[Y_n] = O(n^3)$.

Proof (by induction): will show $E[Y_n] \leq cn^3$


$$\begin{aligned}
 E[Y_n] &\leq \frac{4}{n} \sum_{k=0}^{n-1} E[Y_k] \\
 &\leq \frac{4}{n} \sum_{k=0}^{n-1} ck^3 \\
 &= \frac{4c}{n} \sum_{k=0}^{n-1} k^3 \\
 &\leq \frac{4c}{n} \int_0^n x^3 dx \\
 &= cn^3
 \end{aligned}$$




23

Jensen's Inequality

Definition. A real valued function f defined on an interval is **convex** if for any two points x and y in its domain and $0 \leq t \leq 1$, $f(tx + (1-t)y) \leq tf(x) + (1-t)f(y)$



Theorem. Let X be a random variable and $f: R \rightarrow R$, a convex function. Then $f(E[X]) \leq E[f(X)]$



24

Conclusion

$f(x) = 2^x$ is convex!

$$f(E(X_n)) = 2^{E(X_n)} \leq E(f(X_n)) = E(2^{X_n}) = E(Y_n)$$

$$2^{E(X_n)} \leq E(Y_n) \leq cn^3$$

$$E(X_n) \leq 3 \log n + c'$$


Theorem. The expected height of a binary search tree randomly built on n keys is $O(\log n)$

25

Treap = (Search) tree + Heap

- A treap is a binary tree.
- Each node contains an element x with $\text{key}(x) \in U$ and priority $\rho(x) \in R$.
- The following hold
 - *Search tree property.* Same as in regular BSTs
 - *Heap property.* For all x , $\rho(x) > \rho(\text{parent}[x])$

key	a	b	c	d	e	f	g
priority	3	7	4	1	5	2	6



26

Treap Uniqueness

- **Lemma.** Let X be a set of keys with distinct priorities $\rho : X \rightarrow R$. Then there is a unique binary search tree for X that satisfies the heap order given by ρ .

Proof. By induction

- Structurally, the treap has the structure that would result if the elements were inserted in priority order

27

Randomized Search Trees

- Treaps with random priorities on a subset S of a universe $(U, <)$ of keys with a total order
- Priorities interpreted as “arrival times”
- Operations
 - Search(x, S): Is $x \in S$?
 - Insert(x, S): Insert x into S if not already in S
 - Delete(x, S): Delete x from S
 - Minimum(S): Return smallest key.
 - Maximum(S): Return largest key.
 - Union(S_1, S_2): Merge S_1 and S_2 .
Precondition: $\forall x_1 \in S_1, x_2 \in S_2: x_1 < x_2$
 - Split(S, x, S_1, S_2): Split S into $S_1, \{x\}$ and S_2 .
 $\forall x_1 \in S_1, x_2 \in S_2: x_1 < x$ and $x < x_2$

28


Operations

- Search (x, S) , $\text{Min}(S)$, $\text{Max}(S)$, $\text{Pred}(x, S)$, $\text{Succ}(x, S)$:
 - Same as BSTs \Rightarrow can do in $O(\log n)$ expected time
- $\text{Insert}(x, S)$
- $\text{Delete}(x, S)$
- $\text{Split}(S, x, S_1, S_2)$: Split S into S_1 , $\{x\}$ and S_2 .
 - Goal: $\forall x_1 \in S_1, x_2 \in S_2: x_1 < x \text{ and } x < x_2$
- $\text{Union}(S_1, S_2)$: Merge S_1 and S_2 .
 - Precondition: $\forall x_1 \in S_1, x_2 \in S_2: x_1 < x_2$

29

Rotations


- Basic tree restructuring operation
- Preserves BST (order) property
- Accomplished by $O(1)$ pointer changes



- In both cases, in-order traversal yields
 $\langle \alpha \rangle x \langle \beta \rangle y \langle \gamma \rangle$

30

Example




Operations...

- **Insert(x, S)**
Insert x as in regular BSTs. Assign random $\rho(x)$ in $[0, 1]$. Rotate x up until $\rho(x) > \rho(p[x])$
- **Delete(x, S)**
Change $\rho(x)$ to ∞ . Rotate x down (by rotating child with smaller ρ up) until heap order is restored. Remove x which is now a leaf.
- **Split(S, x, S_1, S_2):** Split S into S_1 , $\{x\}$ and S_2 with $y < x$ if y in S_1 and $y > x$ if y in S_2
Change $\rho(x)$ to $-\infty$. Rotate x up to root. Return $S_1 = \text{left}(x)$ and $S_2 = \text{right}(x)$

32

Red-Black Trees

- A type of binary search trees
 - Additional color attribute: red or black
 - Tree is full (all nodes, except leaves, have degree 2)
 - All leaves are empty (keys reside in internal nodes)
- Balanced: height is $O(\log n)$
- Operations take $\Theta(\log n)$ in worst case




33

RB-Tree Properties

1. Every node is red or black.
2. The root is black.
3. Every leaf ($nil[T]$) is black.
4. If a node is red, then its parent is black.
5. All paths from a node x to descendant leaves contain the same number of black nodes, called the **black height** of x (exclude the color of x when counting).

34


Example



35

Exercise


- Can you turn the following into a RB tree with keys $\langle 3, 7, 8, 10, 12, 18, 23, 27 \rangle$



36

RB-Tree or 2-3-4 Tree?

- Merge each red node into black parent \Rightarrow 2-3-4 Tree
 - Every internal node has 2 to 4 descendants
 - Every leaf has same depth
 - How many leaves for n keys?
- What is the depth of leaves in a 2-3-4 tree?
- What does height of 2-3-4 tree tells us about RB-tree?



Property Consequences

Claim. A node with height h has black height $\geq h/2$

Proof. At most $h/2$ red nodes \Rightarrow at least $h/2$ black nodes

Claim. Subtree rooted at x contains $\geq 2^{bh(x)} - 1$ keys

Proof (by induction). Induction on what? h, n, bh ?

Claim. A RB-tree with n keys has height $\leq 2\log(n+1)$

Proof. $n \geq 2^b - 1 \geq 2^{h/2} - 1$, where h =height, b =black height

Claim. In a RB-tree with n keys Min, Max, Pred, Succ, and Search run in $O(\log n)$ time.


Tree Modifying Operations

- Can proceed as with regular BST, but...
- When inserting what color is the new node?
 - If red \Rightarrow may violate Property 4 (red node \Rightarrow black parent)
 - If black \Rightarrow may violate Property 5 (equal black height)
- When deleting, we remove one node
 - If this node is red we are ok
 - If node is black \Rightarrow may violate properties 2, 4, 5

39

Rotations

- Basic tree restructuring operation
- Preserves BST (order) property
- Accomplished by $O(1)$ pointer changes




- In both cases, in-order traversal yields
 $\langle \alpha \rangle x \langle \beta \rangle y \langle \gamma \rangle$

40

Insert


- Use standard BST insert algorithm
- Initially paint new node z red
- If parent of z is black we are done
else, there are 3 cases, depending on relative position of z wrt $p[z]$, and color of y (z 's uncle)



41

Case 1


- z 's uncle (y) is red
- Relative position of z not relevant



42

Case 2


- y is black
- z 's parent and z are on opposite sides of their parents
- Convert to Case 3



43

Case 3


- y is black
- Both z 's parent and z are left or both right children
- Make $p[z]$ black and $p[p[z]]$ red, then rotate at $p[p[z]]$



44

Example

- Insert 5 into the following tree



45

Analysis

- Regular BST insert now takes only $O(\log n)$ time (why?)
- Each transition in state diagram takes $O(1)$ time and terminates or moves z two levels up
- $O(\log n)$ levels \Rightarrow insert takes $O(\log n)$ time
- Each insert requires 0, 1, or 2 rotations


46

AVL Trees

- Type of “balanced” BST
- Named after Adelson-Velsky and Landis
- for each node, define its **balance factor** as:
height of left subtree – height of right subtree
- balance factor of every node is -1 , 0 , or 1
 - Note: height of *nil* is -1

47

Example: Balance Factors



48

Insertion Algorithm

1. Insert node x using regular BST insert
2. Restructure the tree if necessary
 - Only nodes on the insertion path can have their balance factor altered
 - Walk up the path towards root and update b.f.'s
 - Stop at deepest node α (if any) whose balance factor is out of range
 - Correct imbalance at α via rotations (4 cases)
 - After “fixing” α there is no need to continue up

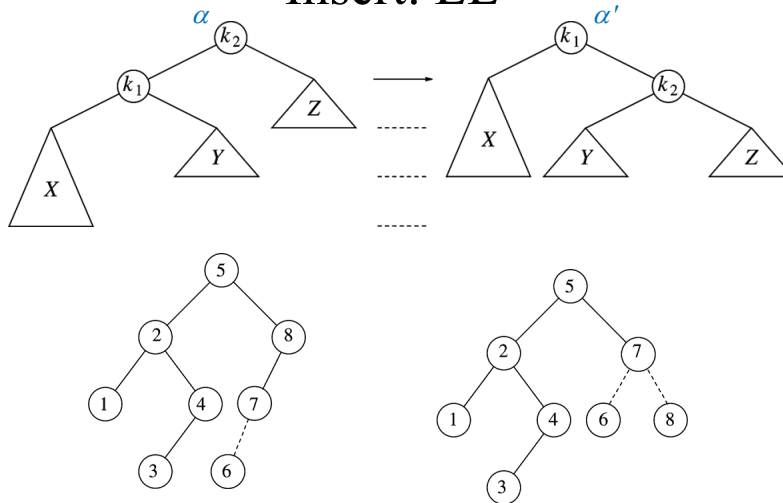
49

Rebalancing

- Let α be the node that needs rebalancing
- The 2 subtrees of α differ by 2 in height
- There are 4 cases, depending on where the insertion of the new value occurred:
 1. In the left subtree of the left child of α (**LL**)
 2. In the right subtree of the left child of α (**LR**)
 3. In the left subtree of the right child of α (**RL**)
 4. In the right subtree of the right child of α (**RR**)

50

Insert: LL




Exercise: Insert 3,2,1,4,5,6,7 into empty tree


51

Insert: LR

- Single rotation does not help



- Can fix with 2 rotations



52

Height of AVL Tree

- $\lfloor \log n \rfloor \leq h < 1.44 \log (n+1) + c$
- $N_h = \text{min \# of nodes in AVL tree of height } h$


$$N_{-1} = 0, N_0 = 1, N_1 = 2, \dots$$

$$N_h = N_{h-1} + N_{h-2} + 1$$

0, 1, 2, 4, 7, 12, 20, 33, 0, 1, 1, 2, 3, 5, 8, 13, 21, 34,

$$F_i = \frac{\varphi^i - (1-\varphi)^i}{\sqrt{5}}, \text{ where } \varphi = \frac{1+\sqrt{5}}{2} \approx 1.6180339887\dots$$

Smallest AVL Tree of Height 9



54